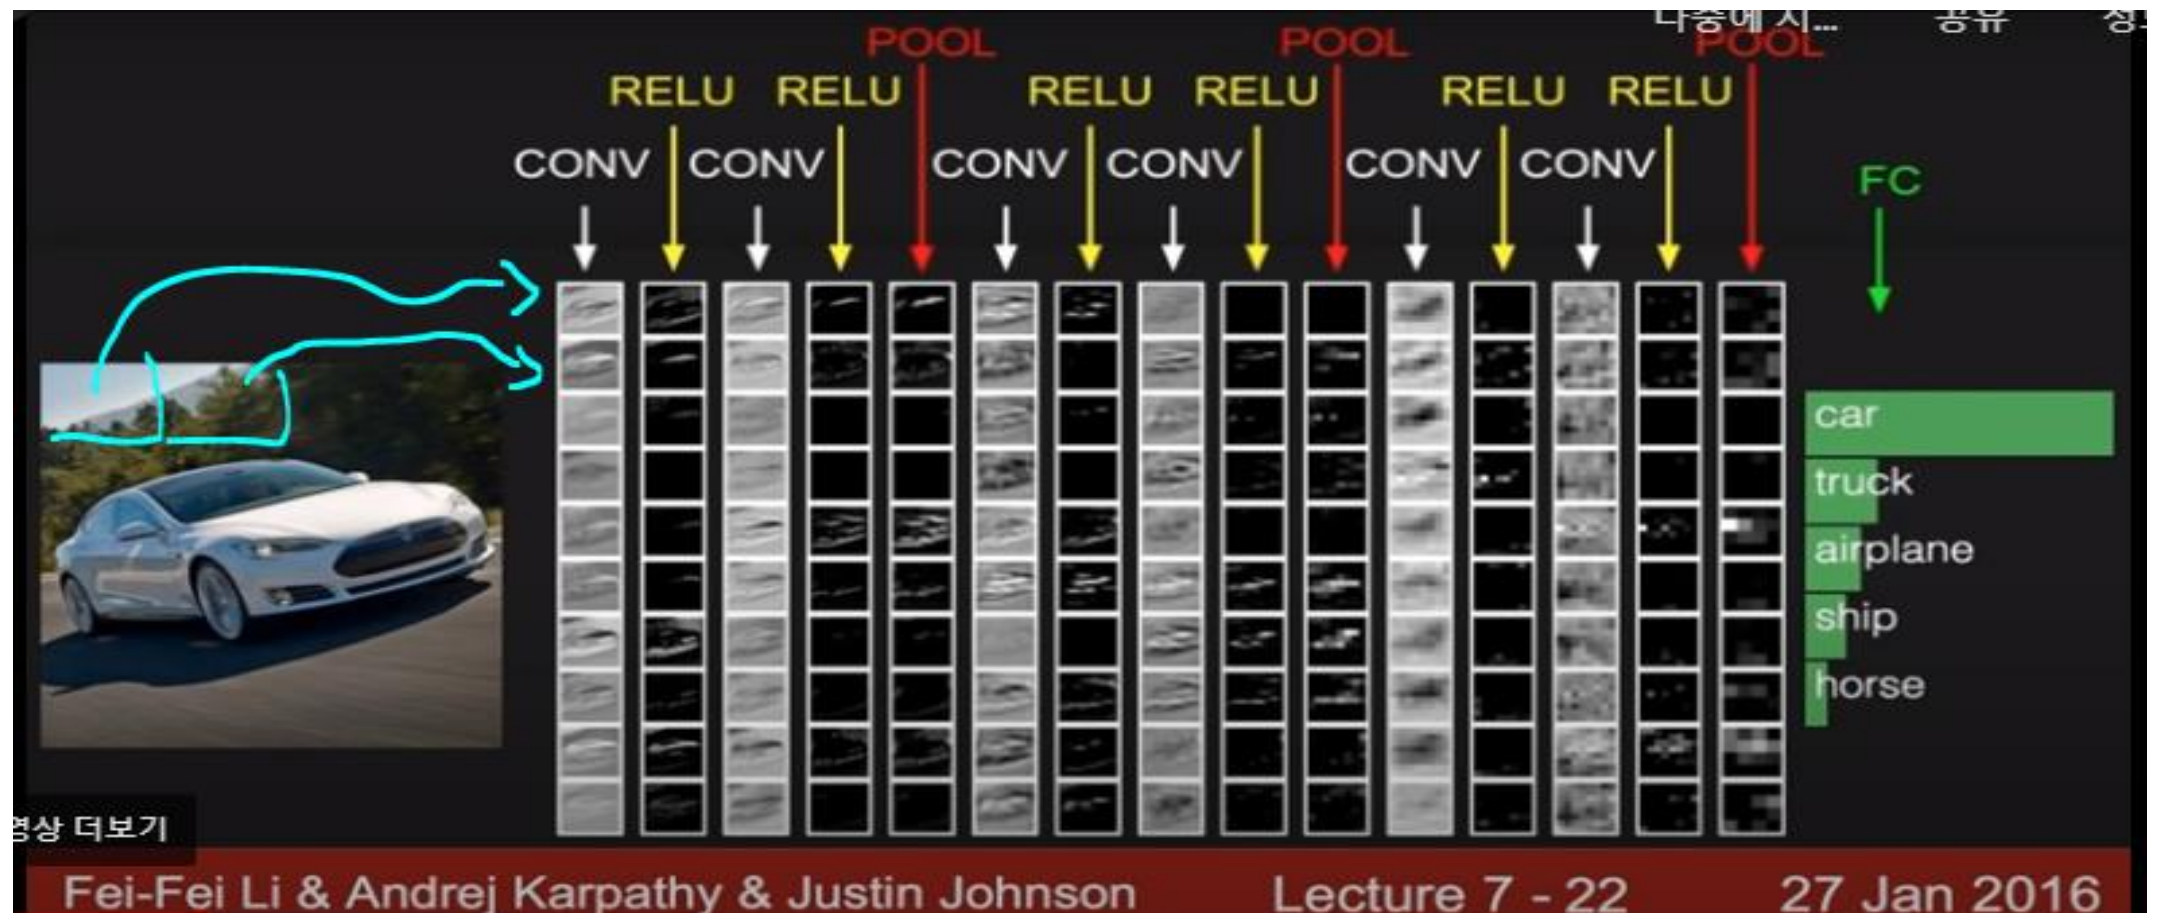




Convolutional Neural Networks

전체구조






Input image (R,G,B)



Input image의 일부분 filter (R,G,B)
Filter 크기 조정 가능



끝까지 filter 단위로 one number들을 뽑아낸다.



Stride = 1
(한 칸씩 움직이는 것)

filter = 3x3

∴ output = 5x5



일반화


Input img : $N \times N$, filter : $F \times F$.

Output : $(N - F) / \text{stride} + 1$

→ 이따가 계속 작아짐 & 문제가 발생하기.



"padding" → 끝에 0으로 채워준다.

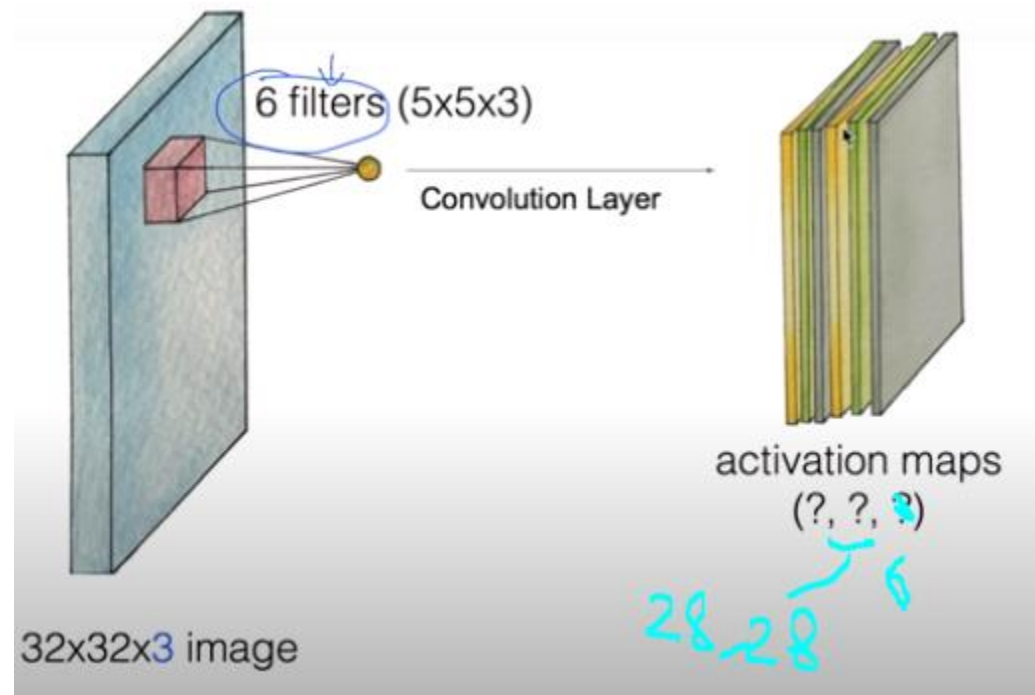



input
→ 9×9

filter
→ 3×3


Output
→ 7×7

볼래 이미지와 같은 size.




여러 개의 필터를 적용해 만든 activation maps
 Input: 32x32, filter:5x5 이므로
 Output : $(32-5)/1+1=28$

Convolution layers



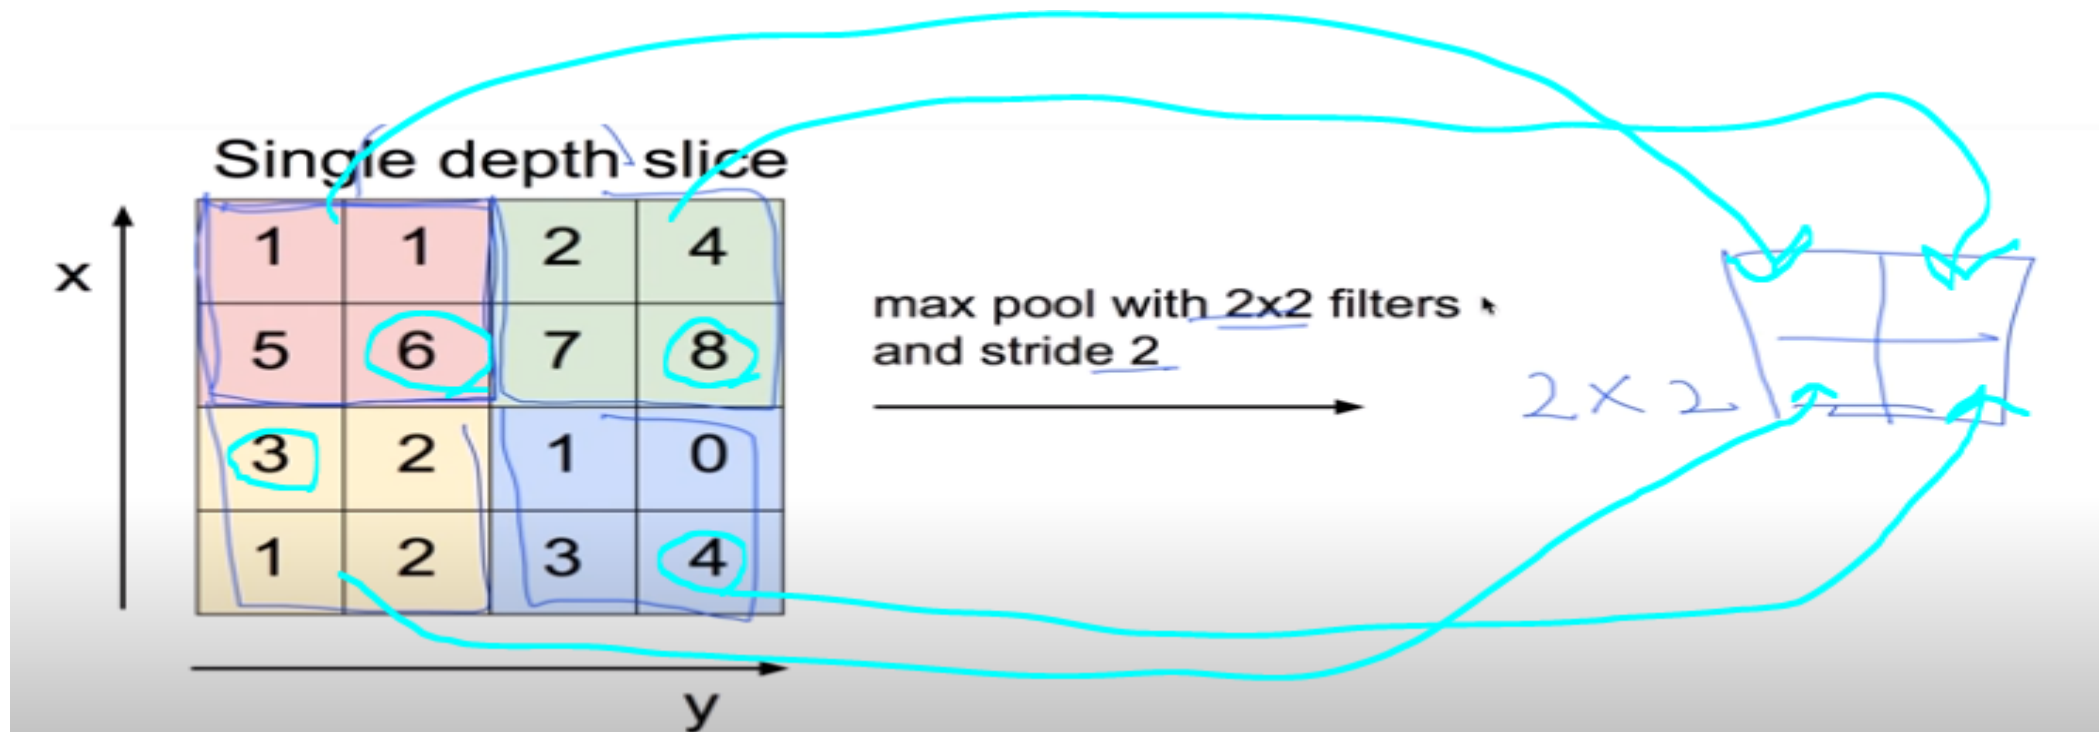
e.g. = filter의 개수

Pooling layer(sampling)



Conv layer에서 1개씩 뽑아내서 resize(sampling)


MAX POOLING



Filter에서 가장 큰 값만 뽑는다.

Case Study: LeNet-5


[LeCun et al., 1998]



Conv filters were 5x5, applied at stride 1
Subsampling (Pooling) layers were 2x2 applied at stride 2
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

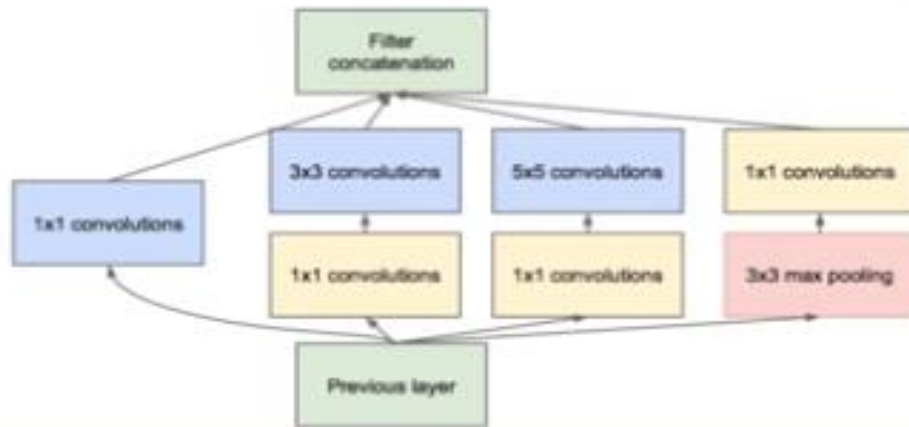

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)

최근에는 normalization 안함

Case Study: GoogLeNet

[Szegedy et al., 2014]



Inception module

ILSVRC 2014 winner (6.7% top 5 error)

Case Study: ResNet

[He et al., 2015]

ILSVRC 2015 winner (3.6% top 5 error)

Research

MSRA @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**
 - ImageNet Classification: *"Ultra-deep"* (quote Yann) **152-layer nets**
 - ImageNet Detection: **16%** better than 2nd
 - ImageNet Localization: **27%** better than 2nd
 - COCO Detection: **11%** better than 2nd
 - COCO Segmentation: **12%** better than 2nd

*improvements are relative numbers

ICCV 15

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.


Slide from Kaiming He's recent presentation <https://www.youtube.com/watch?v=1PGLj-uKT1w>

Fei-Fei Li & Andrej Karpathy & Justin Johnson


Lecture 7 - 77

27 Jan 2016

• Plain net



• Residual net




인간(에러율 5%)보다
좋은 성능

CNN 기본 - input

```
sess = tf.InteractiveSession()
image = np.array([[[[1], [2], [3]],
                   [[4], [5], [6]],
                   [[7], [8], [9]]]], dtype=np.float32)
print(image.shape)
plt.imshow(image.reshape(3,3), cmap='Greys')
```

(1, 3, 3, 1)


<matplotlib.image.AxesImage at 0x1a503dcc948>



Input : N개 이미지, 사이즈, 색상
(1,3,3,1) 1개 img, 3x3, 1(흑백)

Simple convolution layer

Image: 1,3,3,1 image, Filter: 2,2,1,1, Stride: 1x1, Padding: VALID




Weight(=filter)

CNN 기본 – filter , output

```
# print("img: %n", image)
print("image.shape", image.shape)
weight = tf.constant([[[[1.]], [[1.]]],
                      [[1.]], [[1.]]])
print("weight.shape", weight.shape)
conv2d = tf.nn.conv2d(image, weight, strides=[1, 1, 1, 1], padding='VALID')
conv2d_img = conv2d.eval()
print("conv2d_img.shape", conv2d_img.shape)
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    print(one_img.reshape(2,2))
    plt.subplot(1,2,i+1), plt.imshow(one_img.reshape(2,2), cmap='gray')
```

```
image.shape (1, 3, 3, 1)
weight.shape (2, 2, 1, 1)
conv2d_img.shape (1, 2, 2, 1)
[[ 12.  16.]
 [ 24.  28.]
```



Filter(=weight) : 사이즈, 색상, N개 이미지
(2,2,1,1) 2x2, 1(흑백), 1개 img


Output : N개 이미지, 사이즈, 색상
(1,2,2,1) 1개 img, 2x2, 1(흑백)

CNN 기본 – padding(output)

```
# print("imag:\n", image)
print("image.shape", image.shape)

weight = tf.constant([[[[1.]], [[1.]]],
                      [[1.]], [[1.]]])
print("weight.shape", weight.shape)
conv2d = tf.nn.conv2d(image, weight, strides=[1, 1, 1, 1], padding='SAME')
conv2d_img = conv2d.eval()
print("conv2d_img.shape", conv2d_img.shape)
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    print(one_img.reshape(3,3))
    plt.subplot(1,2,i+1), plt.imshow(one_img.reshape(3,3), cmap='gray')
```

```
image.shape (1, 3, 3, 1)
weight.shape (2, 2, 1, 1)
conv2d_img.shape (1, 3, 3, 1)
[[12. 16.  9.]
 [24. 28. 15.]
 [15. 17.  9.]]
```




padding = 'SAME' 을 통해 input이미지와 output 이미지 사이즈가 같게 된다.

CNN 기본 - 필터 3장

```
# print("img:₩n", image)
print("image.shape", image.shape)

weight = tf.constant([[[[1., 10., -1.]], [[1., 10., -1.]]],
                      [[[1., 10., -1.]], [[1., 10., -1.]]]])
print("weight.shape", weight.shape)
conv2d = tf.nn.conv2d(image, weight, strides=[1, 1, 1, 1], padding='SAME')
conv2d_img = conv2d.eval()
print("conv2d_img.shape", conv2d_img.shape)
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    print(one_img.reshape(3,3))
    plt.subplot(1,3,i+1), plt.imshow(one_img.reshape(3,3), cmap='gray')
```

```
image.shape (1, 3, 3, 1)
weight.shape (2, 2, 1, 3)
conv2d_img.shape (1, 3, 3, 3)
[[12. 16.  9.]
 [24. 28. 15.]
 [15. 17.  9.]]
[[120. 160.  90.]
 [240. 280. 150.]
 [150. 170.  90.]]
[[-12. -16.  -9.]
 [-24. -28. -15.]
 [-15. -17.  -9.]]
```



Filter(=weight) : 사이즈, 색상, N개 이미지
(2,2,1,3) 2x2, 1(흑백), 3개 img

CNN 기본 – Max Pooling

4	3
2	1

```
image = np.array([[[[4],[3]],  
                  [[2],[1]]]], dtype=np.float32)  
pool = tf.nn.max_pool(image, ksize=[1, 2, 2, 1],  
                      strides=[1, 1, 1, 1], padding='VALID')  
print(pool.shape)  
print(pool.eval())  
  
(1, 1, 1, 1)  
[[[4.]]]
```

SAME: Zero paddings

4	3	0
2	1	0
0	0	0

4	3	0
2	1	0
0	0	0

4	3	0
2	1	0
0	0	0

4	3	0
2	1	0
0	0	0

```
image = np.array([[[[4],[3]],  
                  [[2],[1]]]], dtype=np.float32)  
pool = tf.nn.max_pool(image, ksize=[1, 2, 2, 1],  
                      strides=[1, 1, 1, 1], padding='SAME')  
print(pool.shape)  
print(pool.eval())  
  
(1, 2, 2, 1)  
[[[4.]  
  [3.]  
  [2.]  
  [1.]]]
```


Output(=ksize) : img에 패딩한 것을 2x2사이즈로 stride 만큼 보면서
max_pool(최대값)반환

padding = 'SAME'하면
패딩 적용한 결과

MNIST

```
img = mnist.train.images[0].reshape(28,28)  
plt.imshow(img, cmap='gray')
```

<matplotlib.image.AxesImage at 0x1a507dceec8>



MNIST Convolution layer


```
sess = tf.InteractiveSession()

img = img.reshape(-1, 28, 28, 1)
W1 = tf.Variable(tf.random_normal([3, 3, 1, 5], stddev=0.01))
conv2d = tf.nn.conv2d(img, W1, strides=[1, 2, 2, 1], padding='SAME')
print(conv2d)
sess.run(tf.global_variables_initializer())
conv2d_img = conv2d.eval()
conv2d_img = np.swapaxes(conv2d_img, 0, 3)
for i, one_img in enumerate(conv2d_img):
    plt.subplot(1, 5, i+1), plt.imshow(one_img.reshape(14, 14), cmap='gray')
```

C:\Users\User\Anaconda3\lib\site-packages\tensorflow\python\client\session.py:1735:
UserWarning: An interactive session is already active. This can cause out-of-memory
errors in some cases. You must explicitly call `InteractiveSession.close()` to relea
se resources held by the other session(s).

warnings.warn('An interactive session is already active. This can ')

Tensor("Conv2D_3:0", shape=(1, 14, 14, 5), dtype=float32)



Filter(=weight) : 사이즈, 색상, N개 이미지
(3,3,1,5) 3x3, 1(흑백), 5개 img


Stride = 2x2

$(28-3)/2+1=14$

MNIST Max Pooling


```
pool = tf.nn.max_pool(conv2d, ksize=[1, 2, 2, 1], strides=[
    1, 2, 2, 1], padding='SAME')
print(pool)
sess.run(tf.global_variables_initializer())
pool_img = pool.eval()
pool_img = np.swapaxes(pool_img, 0, 3)
for i, one_img in enumerate(pool_img):
    plt.subplot(1,5,i+1), plt.imshow(one_img.reshape(7, 7), cmap='gray')
```

Tensor("MaxPool_2:0", shape=(1, 7, 7, 5), dtype=float32)



Output(=ksize) : img에 패딩한 것을 2x2사이즈로 stride 만큼 보면서 max_pool(최대값)반환

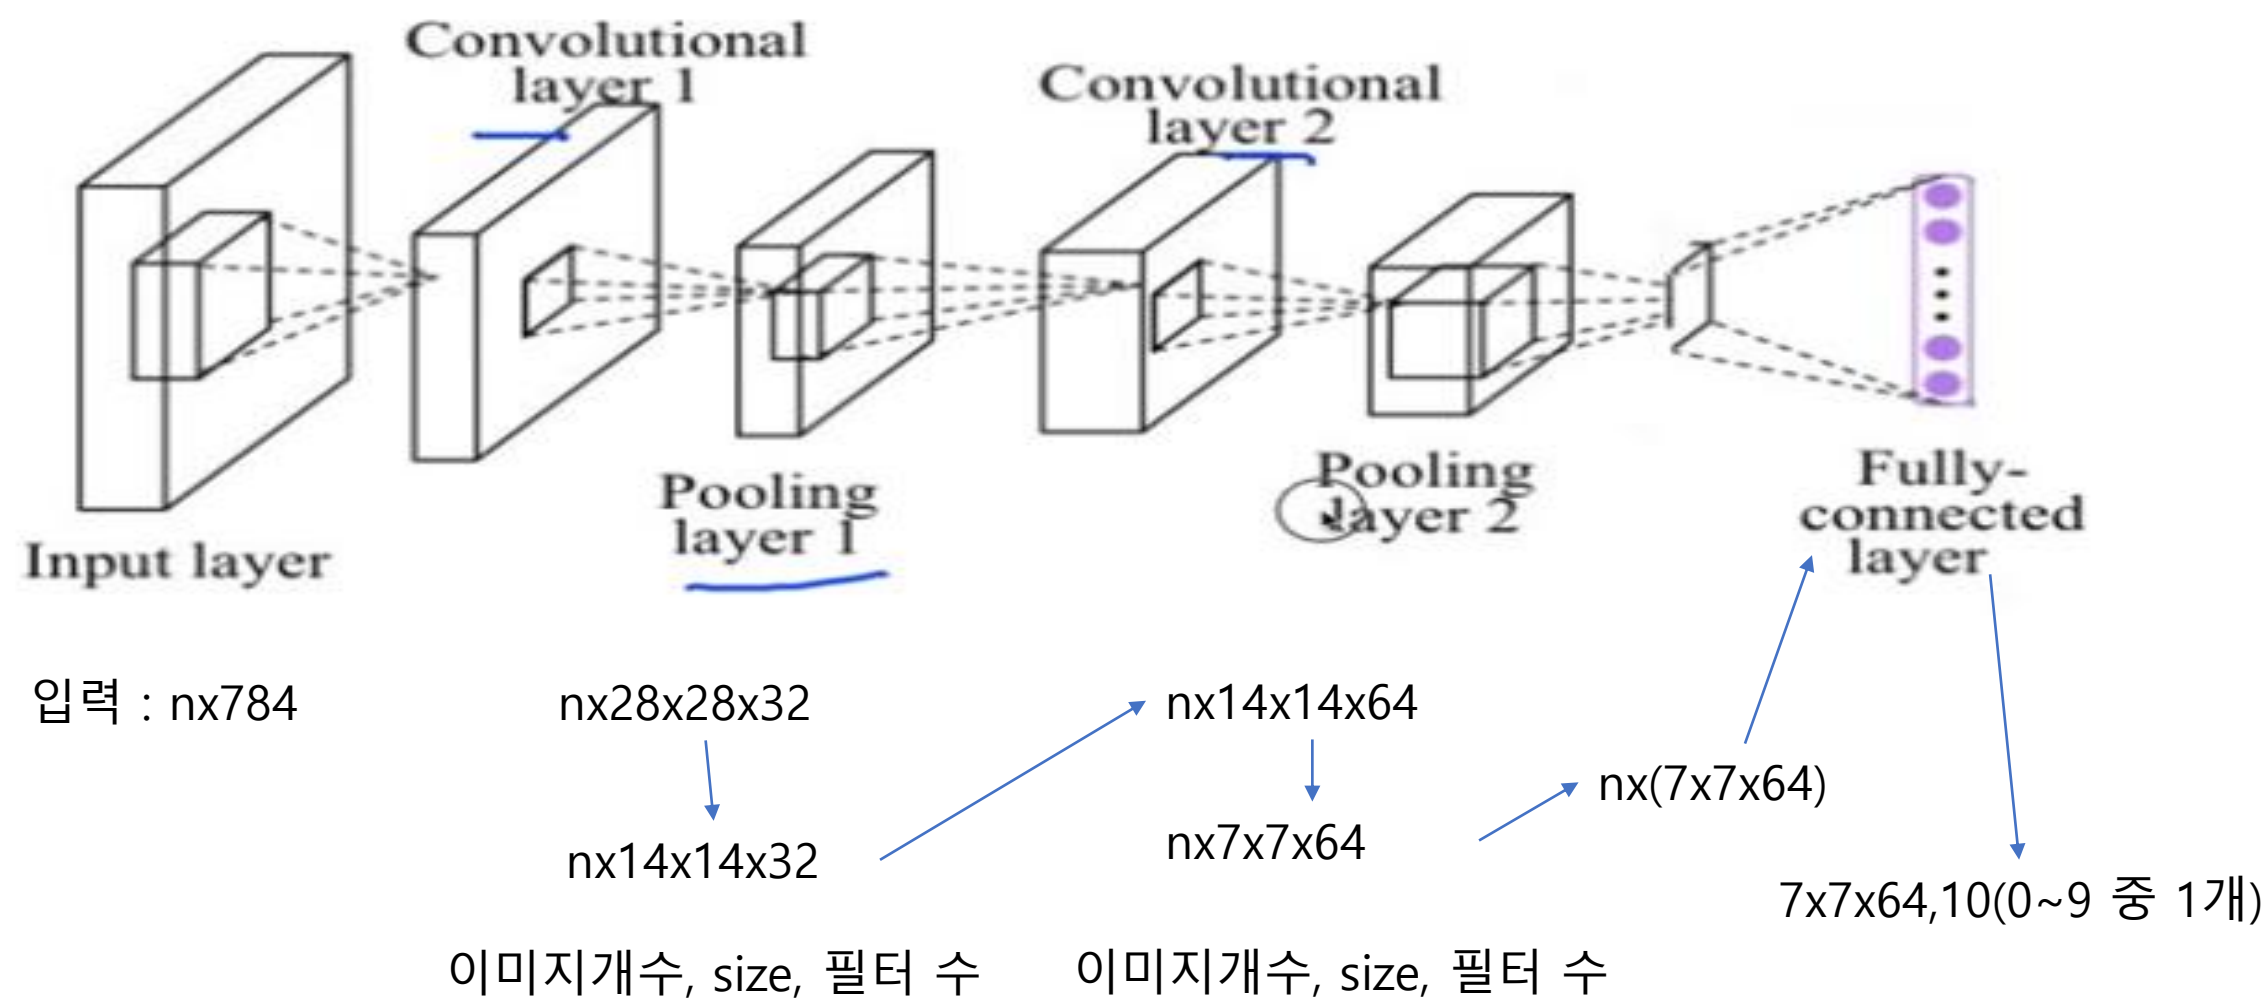
CNN MNIST: 99%



이미지를 받아 컨벌루션 통과 후 subsampling 반복 후 fully connected에 연결한다.

Simple CNN

98% 정확도




```

# 입력 nx784
X = tf.placeholder(tf.float32, [None, 784])
# 맨 앞 : n개(자동으로 맞춰줄) 이미지, 28x28x1 (black/white)
X_img = tf.reshape(X, [-1, 28, 28, 1])
# 나오는 결과
Y = tf.placeholder(tf.float32, [None, 10])

```

입력 이미지 포맷 : n개, size(28x28), channel(1, grayscale)

```

# L1 img in shape=(?, 28, 28, 1)
# 필터 : 3x3크기, 채널 1개, 필터 개수 : 32개
W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
# Conv -> (?, 28, 28, 32)
# Pool -> (?, 14, 14, 32)
# padding 하고 stride가 1x1이므로 입력 이미지와 결과 이미지의 size가 동일하다.
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
L1 = tf.nn.relu(L1)
# stride가 2x2이므로 14x14 이미지가 나온다. (L1결과)
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
...
Tensor("Conv2D:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("Relu:0", shape=(?, 28, 28, 32), dtype=float32)
Tensor("MaxPool:0", shape=(?, 14, 14, 32), dtype=float32)
...

```

L1 : nxsize(28x28)x32 \longrightarrow L1 : nxsize(14x14)x32

```

# L1에서 나온 결과가 L2의 입력으로 들어간다.
# L2 img in shape=(?, 14, 14, 32)
# 필터 : 3x3, 채널 32개, 필터의 개수 : 64개
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))
# Conv -> (?, 14, 14, 64)
# Pool -> (?, 7, 7, 64)
L2 = tf.nn.conv2d(L1, W2, strides=[1, 1, 1, 1], padding='SAME')
L2 = tf.nn.relu(L2)
# stride 2x2해서 이미지의 사이즈가 절반으로 된다. (7x7)
L2 = tf.nn.max_pool(L2, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
# n개들을 7x7x64로 펼쳐준다.
L2_flat = tf.reshape(L2, [-1, 7 * 7 * 64])
...

Tensor("Conv2D_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:0", shape=(?, 7, 7, 64), dtype=float32)
Tensor("Reshape_1:0", shape=(?, 3136), dtype=float32)
...

```

L2 : nxsize(28x28)x64 \longrightarrow L2 : nxsize(14x14)x64

```

# L2에서 나온 결과(L2_flat)를 마지막 fully-connected layer에 넣어준다.
# Final FC 7x7x64 inputs -> 10 outputs
W3 = tf.get_variable("W3", shape=[7 * 7 * 64, 10],
                    initializer=tf.contrib.layers.xavier_initializer())
b = tf.Variable(tf.random_normal([10]))
logits = tf.matmul(L2_flat, W3) + b

```

L2_flat : nx(7*7*64) \longrightarrow F: (7*7*64)x10

```
# define cost/loss & optimizer
```

```
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(  
    logits=logits, labels=Y))
```

```
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
```

```
# initialize
```

```
sess = tf.Session()
```

```
sess.run(tf.global_variables_initializer())
```

Cost function 정의,
session 열기

```
# train my model
```

```
print('Learning started. It takes sometime.')
```

```
for epoch in range(training_epochs):
```

```
    avg_cost = 0
```

```
    total_batch = int(mnist.train.num_examples / batch_size)
```

```
    for i in range(total_batch):
```

```
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)
```

```
        feed_dict = {X: batch_xs, Y: batch_ys}
```

```
        c, _ = sess.run([cost, optimizer], feed_dict=feed_dict)
```

```
        avg_cost += c / total_batch
```

```
    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))
```

```
print('Learning Finished!')
```

모델 학습시키기

```
# Test model and check accuracy
```

```
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(Y, 1))
```

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

```
print('Accuracy:', sess.run(accuracy, feed_dict={  
    X: mnist.test.images, Y: mnist.test.labels}))
```

```
# Get one and predict
```

```
r = random.randint(0, mnist.test.num_examples - 1)
```

```
print("Label: ", sess.run(tf.argmax(mnist.test.labels[r:r + 1], 1)))
```

```
print("Prediction: ", sess.run(  
    tf.argmax(logits, 1), feed_dict={X: mnist.test.images[r:r + 1]}))
```

```
# plt.imshow(mnist.test.images[r:r + 1].
```


```
#             reshape(28, 28), cmap='Greys', interpolation='nearest')
```

```
# plt.show()
```

Test(정확도) 측정,
Prediction

Deep CNN

99% 정확도



```

# L3 Input shape=(?, 7, 7, 64)
W3 = tf.Variable(tf.random_normal([3, 3, 64, 128], stddev=0.01))
# Conv ->(?, 7, 7, 128)
# Pool ->(?, 4, 4, 128)
# Reshape ->(?, 4 * 4 * 128) # Flatten them for FC
L3 = tf.nn.conv2d(L2, W3, strides=[1, 1, 1, 1], padding='SAME')
L3 = tf.nn.relu(L3)
L3 = tf.nn.max_pool(L3, ksize=[1, 2, 2, 1], strides=[
    1, 2, 2, 1], padding='SAME')
L3 = tf.nn.dropout(L3, keep_prob=keep_prob)
L3_flat = tf.reshape(L3, [-1, 128 * 4 * 4])

Tensor("Conv2D_2:0", shape=(?, 7, 7, 128), dtype=float32)
Tensor("Relu_2:0", shape=(?, 7, 7, 128), dtype=float32)
Tensor("MaxPool_2:0", shape=(?, 4, 4, 128), dtype=float32)
Tensor("dropout_2/mul:0", shape=(?, 4, 4, 128), dtype=float32)
Tensor("Reshape_1:0", shape=(?, 2048), dtype=float32)

```

L3 : nxsize(7x7)x128 \longrightarrow L3 : nxsize(4x4)x128

```

# L4 FC 4x4x128 inputs -> 625 outputs
W4 = tf.get_variable("W4", shape=[128 * 4 * 4, 625],
    initializer=tf.contrib.layers.xavier_initializer())
b4 = tf.Variable(tf.random_normal([625]))
L4 = tf.nn.relu(tf.matmul(L3_flat, W4) + b4)
L4 = tf.nn.dropout(L4, keep_prob=keep_prob)

Tensor("Relu_3:0", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul:0", shape=(?, 625), dtype=float32)

```

L3_flat : nx(4*4*128) \longrightarrow F1: (4*4*128)x625

```
# L5 Final FC 625 inputs -> 10 outputs
W5 = tf.get_variable("W5", shape=[625, 10],
                      initializer=tf.contrib.layers.xavier_initializer())
b5 = tf.Variable(tf.random_normal([10]))
logits = tf.matmul(L4, W5) + b5
...
```

Tensor("add_1:0", shape=(?, 10), dtype=float32)
...

F1 : nx625



F2: 625x10

CNN_class

```
class Model:|  
    def __init__(self, sess, name):  
        self.sess = sess  
        self.name = name  
        self._build_net()
```

```
def _build_net(self):  
    with tf.variable_scope(self.name):  
        # dropout (keep_prob) rate 0.7~0.5 on training, but should be 1  
        # for testing  
        self.keep_prob = tf.placeholder(tf.float32)  
  
        # input place holders  
        self.X = tf.placeholder(tf.float32, [None, 784])  
        # img 28x28x1 (black/white)  
        X_img = tf.reshape(self.X, [-1, 28, 28, 1])  
        self.Y = tf.placeholder(tf.float32, [None, 10])  
  
        # L1 imgIn shape=(?, 28, 28, 1)  
        W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))  
        # Conv -> (?, 28, 28, 32)  
        # Pool -> (?, 14, 14, 32)  
        L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')  
        L1 = tf.nn.relu(L1)  
        L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],  
                             strides=[1, 2, 2, 1], padding='SAME')  
        L1 = tf.nn.dropout(L1, keep_prob=self.keep_prob)  
        ...  
  
        Tensor("Conv2D:0", shape=(?, 28, 28, 32), dtype=float32)  
        Tensor("Relu:0", shape=(?, 28, 28, 32), dtype=float32)  
        Tensor("MaxPool:0", shape=(?, 14, 14, 32), dtype=float32)  
        Tensor("dropout/mul:0", shape=(?, 14, 14, 32), dtype=float32)  
        ...
```

CNN_class

```
# L2 img in shape=(?, 14, 14, 32)
W2 = tf.Variable(tf.random_normal([3, 3, 32, 64], stddev=0.01))
# Conv ->(?, 14, 14, 64)
# Pool ->(?, 7, 7, 64)
L2 = tf.nn.conv2d(L1, W2, strides=[1, 1, 1, 1], padding='SAME')
L2 = tf.nn.relu(L2)
L2 = tf.nn.max_pool(L2, ksize=[1, 2, 2, 1],
                    strides=[1, 2, 2, 1], padding='SAME')
L2 = tf.nn.dropout(L2, keep_prob=self.keep_prob)
...

Tensor("Conv2D_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("Relu_1:0", shape=(?, 14, 14, 64), dtype=float32)
Tensor("MaxPool_1:0", shape=(?, 7, 7, 64), dtype=float32)
Tensor("dropout_1/mul:0", shape=(?, 7, 7, 64), dtype=float32)
...

# L3 img in shape=(?, 7, 7, 64)
W3 = tf.Variable(tf.random_normal([3, 3, 64, 128], stddev=0.01))
# Conv ->(?, 7, 7, 128)
# Pool ->(?, 4, 4, 128)
# Reshape ->(?, 4 * 4 * 128) # Flatten them for FC
L3 = tf.nn.conv2d(L2, W3, strides=[1, 1, 1, 1], padding='SAME')
L3 = tf.nn.relu(L3)
L3 = tf.nn.max_pool(L3, ksize=[1, 2, 2, 1], strides=[
    1, 2, 2, 1], padding='SAME')
L3 = tf.nn.dropout(L3, keep_prob=self.keep_prob)

L3_flat = tf.reshape(L3, [-1, 128 * 4 * 4])
...

Tensor("Conv2D_2:0", shape=(?, 7, 7, 128), dtype=float32)
Tensor("Relu_2:0", shape=(?, 7, 7, 128), dtype=float32)
Tensor("MaxPool_2:0", shape=(?, 4, 4, 128), dtype=float32)
Tensor("dropout_2/mul:0", shape=(?, 4, 4, 128), dtype=float32)
Tensor("Reshape_1:0", shape=(?, 2048), dtype=float32)
...
```

```
# L4 FC 4x4x128 inputs -> 625 outputs
W4 = tf.get_variable("W4", shape=[128 * 4 * 4, 625],
                    initializer=tf.contrib.layers.xavier_initializer())
b4 = tf.Variable(tf.random_normal([625]))
L4 = tf.nn.relu(tf.matmul(L3_flat, W4) + b4)
L4 = tf.nn.dropout(L4, keep_prob=self.keep_prob)
...

Tensor("Relu_3:0", shape=(?, 625), dtype=float32)
Tensor("dropout_3/mul:0", shape=(?, 625), dtype=float32)
...
```

```
# L5 Final FC 625 inputs -> 10 outputs
W5 = tf.get_variable("W5", shape=[625, 10],
                    initializer=tf.contrib.layers.xavier_initializer())
b5 = tf.Variable(tf.random_normal([10]))
self.logits = tf.matmul(L4, W5) + b5
...

Tensor("add_1:0", shape=(?, 10), dtype=float32)
...
```

```
# define cost/loss & optimizer
self.cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    logits=self.logits, labels=self.Y))
self.optimizer = tf.train.AdamOptimizer(
    learning_rate=learning_rate).minimize(self.cost)

correct_prediction = tf.equal(
    tf.argmax(self.logits, 1), tf.argmax(self.Y, 1))
self.accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```


CNN_class

```
def predict(self, x_test, keep_prop=1.0):  
    return self.sess.run(self.logits, feed_dict={self.X: x_test, self.keep_prob: keep_prop})  
  
def get_accuracy(self, x_test, y_test, keep_prop=1.0):  
    return self.sess.run(self.accuracy, feed_dict={self.X: x_test, self.Y: y_test, self.keep_prob: keep_prop})  
  
def train(self, x_data, y_data, keep_prop=0.7):  
    return self.sess.run([self.cost, self.optimizer], feed_dict={  
        self.X: x_data, self.Y: y_data, self.keep_prob: keep_prop})
```

```
# initialize  
sess = tf.Session()  
m1 = Model(sess, "m1")  
  
sess.run(tf.global_variables_initializer())  
  
print('Learning Started!')
```

```
# train my model  
for epoch in range(training_epochs):  
    avg_cost = 0  
    total_batch = int(mnist.train.num_examples / batch_size)  
  
    for i in range(total_batch):  
        batch_xs, batch_ys = mnist.train.next_batch(batch_size)  
        c, _ = m1.train(batch_xs, batch_ys)  
        avg_cost += c / total_batch  
  
    print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.9f}'.format(avg_cost))  
  
print('Learning Finished!')  
  
# Test model and check accuracy  
print('Accuracy:', m1.get_accuracy(mnist.test.images, mnist.test.labels))
```

CNN_layers

```
def _build_net(self):
    with tf.variable_scope(self.name):
        # dropout (keep_prob) rate 0.7~0.5 on training, but should be 1
        # for testing
        self.keep_prob = tf.placeholder(tf.float32)

        # input place holders
        self.X = tf.placeholder(tf.float32, [None, 784])
        # img 28x28x1 (black/white)
        X_img = tf.reshape(self.X, [-1, 28, 28, 1])
        self.Y = tf.placeholder(tf.float32, [None, 10])

        # L1 img in shape=(?, 28, 28, 1)
        W1 = tf.Variable(tf.random_normal([3, 3, 1, 32], stddev=0.01))
        # Conv -> (?, 28, 28, 32)
        # Pool -> (?, 14, 14, 32)
        L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
        L1 = tf.nn.relu(L1)
        L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1],
                            strides=[1, 2, 2, 1], padding='SAME')
        L1 = tf.nn.dropout(L1, keep_prob=self.keep_prob)
        ...
        Tensor("Conv2D:0", shape=(?, 28, 28, 32), dtype=float32)
        Tensor("Relu:0", shape=(?, 28, 28, 32), dtype=float32)
        Tensor("MaxPool:0", shape=(?, 14, 14, 32), dtype=float32)
        Tensor("dropout/mul:0", shape=(?, 14, 14, 32), dtype=float32)
        ...
```

```
def _build_net(self):
    with tf.variable_scope(self.name):
        # dropout (keep_prob) rate 0.7~0.5 on training, but should be 1
        # for testing
        self.training = tf.placeholder(tf.bool)

        # input place holders
        self.X = tf.placeholder(tf.float32, [None, 784])

        # img 28x28x1 (black/white), Input Layer
        X_img = tf.reshape(self.X, [-1, 28, 28, 1])
        self.Y = tf.placeholder(tf.float32, [None, 10])

        # Convolutional Layer #1
        conv1 = tf.layers.conv2d(inputs=X_img, filters=32, kernel_size=[3, 3],
                                padding="SAME", activation=tf.nn.relu)

        # Pooling Layer #1
        pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2],
                                         padding="SAME", strides=2)

        dropout1 = tf.layers.dropout(inputs=pool1,
                                     rate=0.3, training=self.training)
```

차이점 : conv2d를 쓰면 input을 조건(filter, kernel_size 등)에 적용시킨 결과를 도출한다.

Dropout : 학습인지, Test인지 (rate=1) 구별

CNN_layers

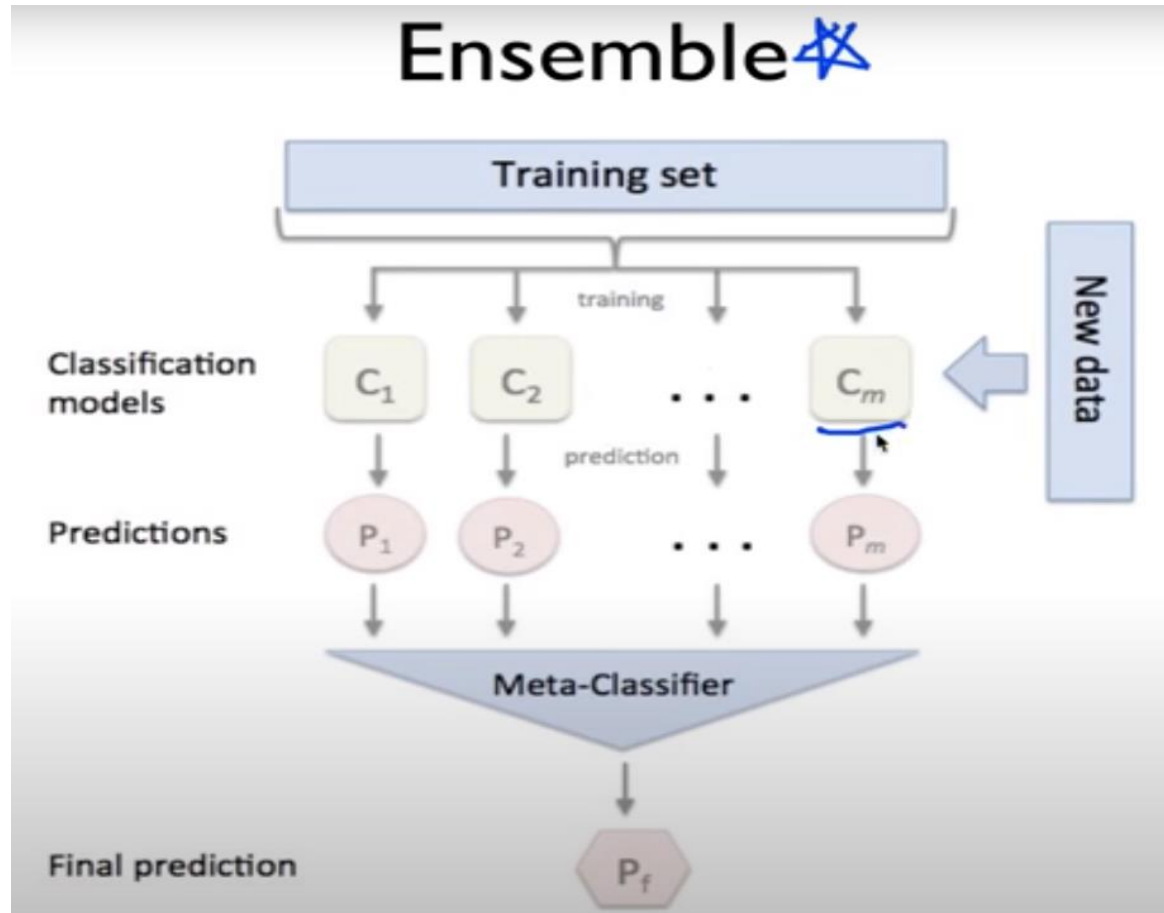
```
# Convolutional Layer #2 and Pooling Layer #2
conv2 = tf.layers.conv2d(inputs=dropout1, filters=64, kernel_size=[3, 3],
                        padding="SAME", activation=tf.nn.relu)
pool2 = tf.layers.max_pooling2d(inputs=conv2, pool_size=[2, 2],
                                padding="SAME", strides=2)
dropout2 = tf.layers.dropout(inputs=pool2,
                             rate=0.3, training=self.training)

# Convolutional Layer #2 and Pooling Layer #2
conv3 = tf.layers.conv2d(inputs=dropout2, filters=128, kernel_size=[3, 3],
                        padding="same", activation=tf.nn.relu)
pool3 = tf.layers.max_pooling2d(inputs=conv3, pool_size=[2, 2],
                                padding="same", strides=2)
dropout3 = tf.layers.dropout(inputs=pool3,
                             rate=0.3, training=self.training)

# Dense Layer with Relu
flat = tf.reshape(dropout3, [-1, 128 * 4 * 4])
dense4 = tf.layers.dense(inputs=flat,
                        units=625, activation=tf.nn.relu)
dropout4 = tf.layers.dropout(inputs=dense4,
                             rate=0.5, training=self.training)
```


Dense: units(출력을 몇 개로 할 것인지), activation을 어떤 것으로 할 것인지만 결정해주면 자동으로 input에 대해 계산해줌
(Fully-connected-layer)

CNN_Ensemble



여러 개의 모델을 training 시키고
새로운 데이터(testing할 데이터)가 들어왔을 때
각 모델들을 예측 시키고
조합한 뒤 최종 결과를 낸다.

CNN_Ensemble



CNN_Ensemble

```
# Test model and check accuracy
test_size = len(mnist.test.labels)
predictions = np.zeros([test_size, 10])
for m_idx, m in enumerate(models):
    print(m_idx, 'Accuracy:', m.get_accuracy(
        mnist.test.images, mnist.test.labels))
    p = m.predict(mnist.test.images)
    predictions += p

ensemble_correct_prediction = tf.equal(
    tf.argmax(predictions, 1), tf.argmax(mnist.test.labels, 1))
ensemble_accuracy = tf.reduce_mean(
    tf.cast(ensemble_correct_prediction, tf.float32))
print('Ensemble accuracy:', sess.run(ensemble_accuracy))
```

여러 개의 모델을 training 시키고
새로운 데이터(testing할 데이터)가 들어왔을 때
각 모델들을 예측 시키고
조합한 뒤 최종 결과를 낸다.
⇒ 각 모델마다 label에 대해 예측한 값들을 다 더한다.
⇒ 최종적으로 다 더한 예측 값들 중 가장 큰(argmax) 값을 선택한다.